

CM146, Fall 2019  
Problem Set 4: Learning Theory, Boosting, Multi-class  
Classification

Due December 8, 2019, 11:59pm

## Submission instructions

- Submit your solutions electronically on the course Gradescope site as PDF files.
- If you plan to typeset your solutions, please use the LaTeX solution template. If you must submit scanned handwritten solutions, please use a black pen on blank white paper and a high-quality scanner app.

## 1 PAC Learning [25 pts]

Let  $\mathcal{X}$  be a set of numbers  $\{1, 2, \dots, N\}$ , and let  $H_{\text{Singleton}} = \{h_z : z \in \mathcal{X}\} \cup \{h^-\}$  is the hypothesis space, where for each  $z \in \mathcal{X}$ ,  $h_z$  is the function defined by  $h_z(x) = 1$  if  $x = z$  and  $h_z(x) = 0$  if  $x \neq z$ .  $h^-$  is simply the all-negative hypothesis, namely,  $\forall x \in X, h^-(x) = 0$ . (For any hypothesis  $h$ , it only receives input  $x \in \mathcal{X}$ , which means the domain is discrete, consisting  $N$  numbers).

Assume a training dataset  $S_{\text{train}}$  consist of  $m$  data points drawn i.i.d from a uniform distribution  $D$  (each data point  $x \in X$  has equal probability, i.e.,  $\frac{1}{N}$ , to be sampled), and that the labels are provided from some target function  $h^* \in H_{\text{Singleton}}$  (true hypothesis belongs to our learning hypothesis family). We simply choose 0-1 loss as train and generalization error metric.

- (a) **[10 points]** Describe an algorithm  $A$  that learns a hypothesis  $A(S_{\text{train}})$  from hypothesis family  $H_{\text{Singleton}}$  on training set  $S_{\text{train}}$ , so that the training error is minimized (such an algorithm is Empirical Risk Minimization (ERM)).
- (b) **[15 points]** Prove that if training set  $S_{\text{train}}$  has a sample size larger than  $\frac{\log(1/\sigma)}{\epsilon}$ , then the probability that the generalization error of  $A(S_{\text{train}})$  is larger than  $\epsilon$  is at most  $\sigma$  (This means that  $H_{\text{Singleton}}$  is PAC learnable), i.e.:

$$\mathbf{P}\left(L_{(D, h^*)}(A(S_{\text{train}})) > \epsilon\right) \leq \sigma \quad (1)$$

where the generalization error is defined as:

$$L_{(D, h^*)}(A(S_{\text{train}})) \quad (2)$$

$$= \mathbf{E}_{x \sim D} \left( 0\text{-}1\text{-Loss}(h^*(x), A(S_{\text{train}})(x)) \right) \quad (3)$$

$$= \mathbf{P}_{x \sim D} \left( h^*(x) \neq A(S_{\text{train}})(x) \right) \quad (4)$$

**Hint:** Write down the generalization error of  $A(S_{\text{train}})$ , estimate the probability that it's larger than  $\epsilon$ , and bound this probability by  $\sigma$ . The following inequality is useful:  $1 - x \leq e^{-x}$ .

$i$	Label	Hypothesis 1 (1st iteration)				Hypothesis 2 (2nd iteration)			
		$D_0$	$f_1 \equiv [x > \_]$	$f_2 \equiv [y > \_]$	$h_1 \equiv [ \_ ]$	$D_1$	$f_1 \equiv [x > \_]$	$f_2 \equiv [y > \_]$	$h_2 \equiv [ \_ ]$
(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
1	−								
2	−								
3	+								
4	−								
5	−								
6	−								
7	+								
8	−								
9	+								
10	+								

Table 1: Table for Boosting results

## 2 VC Dimension [15 pts]

We define the following hypothesis set:

$$H = \{sgn(ax^2 + bx + c); a, b, c, \in R\},$$

where  $sgn(\cdot)$  is 1 when the argument  $\cdot$  is positive, and 0 otherwise. What is the VC dimension of  $H$ ? Prove your claim.

## 3 Boosting [40 pts]

Consider the following examples  $(x, y) \in \mathbb{R}^2$  ( $i$  is the example index):

$i$	$x$	$y$	Label
1	0	8	−
2	1	4	−
3	3	7	+
4	-2	1	−
5	-1	13	−
6	9	11	−
7	12	7	+
8	-7	-1	−
9	-3	12	+
10	5	9	+

In this problem, you will use Boosting to learn a hidden Boolean function from this set of examples. We will use two rounds of AdaBoost to learn a hypothesis for this data set. In each round, AdaBoost chooses a weak learner that minimizes the error  $\epsilon$ . As weak learners, use hypotheses of the form (a)  $f_1 \equiv [x > \theta_x]$  or (b)  $f_2 \equiv [y > \theta_y]$ , for some integers  $\theta_x, \theta_y$  (either one of the two forms, not a

disjunction of the two). There should be no need to try many values of  $\theta_x, \theta_y$ ; appropriate values should be clear from the data. When using log, use base 2.

- (a) [10 points] Start the first round with a uniform distribution  $D_0$ . Place the value for  $D_0$  for each example in the third column of Table 1. Write the new representation of the data in terms of the *rules of thumb*,  $f_1$  and  $f_2$ , in the fourth and fifth columns of Table 1.
- (b) [10 points] Find the hypothesis given by the weak learner that minimizes the error  $\epsilon$  for that distribution. Place this hypothesis as the heading to the sixth column of Table 1, and give its prediction for each example in that column.
- (c) [10 points] Now compute  $D_1$  for each example, find the new best weak learners  $f_1$  and  $f_2$ , and select hypothesis that minimizes error on this distribution, placing these values and predictions in the seventh to tenth columns of Table 1.
- (d) [10 points] Write down the final hypothesis produced by AdaBoost.

**What to submit:** Fill out Table 1 as explained, show computation of  $\alpha$  and  $D_1(i)$ , and give the final hypothesis,  $H_{final}$ .

## 4 Multi-class classification [60 pts]

Consider a multi-class classification problem with  $k$  class labels  $\{1, 2, \dots, k\}$ . Assume that we are given  $m$  examples, labeled with one of the  $k$  class labels. Assume, for simplicity, that we have  $m/k$  examples of each type.

Assume that you have a learning algorithm  $L$  that can be used to learn Boolean functions. (E.g., think about  $L$  as the Perceptron algorithm). We would like to explore several ways to develop learning algorithms for the multi-class classification problem.

There are two schemes to use the algorithm  $L$  on the given data set, and produce a multi-class classification:

- **One vs. All:** For every label  $i \in [1, k]$ , a classifier is learned over the following data set: the examples labeled with the label  $i$  are considered “positive”, and examples labeled with any other class  $j \in [1, k], j \neq i$  are considered “negative”.
- **All vs. All:** For every pair of labels  $\langle i, j \rangle$ , a classifier is learned over the following data set: the examples labeled with one class  $i \in [1, k]$  are considered “positive”, and those labeled with the other class  $j \in [1, k], j \neq i$  are considered “negative”.

- (a) [20 points] For each of these two schemes, answer the following:
  - i. How many classifiers do you learn?
  - ii. How many examples do you use to learn each classifier within the scheme?
  - iii. How will you decide the final class label (from  $\{1, 2, \dots, k\}$ ) for each example?
  - iv. What is the computational complexity of the training process?

- (b) **[5 points]** Based on your analysis above of two schemes individually, which scheme would you prefer? Justify.
- (c) **[5 points]** You could also use a KERNELPERCEPTRON for a two-class classification. We could also use the algorithm to learn a multi-class classification. Does using a KERNELPERCEPTRON change your analysis above? Specifically, what is the computational complexity of using a KERNELPERCEPTRON and which scheme would you prefer when using a KERNELPERCEPTRON?
- (d) **[10 points]** We are given a magical black-box binary classification algorithm (we don't know how it works, but it just does!) which has a learning time complexity of  $O(dn^2)$ , where  $n$  is the total number of training examples supplied (positive+negative) and  $d$  is the dimensionality of each example. What are the overall training time complexities of the all-vs-all and the one-vs-all paradigms, respectively, and which training paradigm is most efficient?
- (e) **[10 points]** We are now given another magical black-box binary classification algorithm (wow!) which has a learning time complexity of  $O(d^2n)$ , where  $n$  is the total number of training examples supplied (positive+negative) and  $d$  is the dimensionality of each example. What are the overall training time complexities of the all-vs-all and the one-vs-all paradigms, respectively, and which training paradigm is most efficient, when using this new classifier?
- (f) **[10 points]** Suppose we have learnt an all-vs-all multi-class classifier and now want to proceed to predicting labels on unseen examples.

We have learnt a simple linear classifier with a weight vector of dimensionality  $d$  for each of the  $k(k-1)/2$  classifier ( $w_i^T x = 0$  is the simple linear classifier hyperplane for each  $i = [1, \dots, k(k-1)/2]$ )

We have two evaluation strategies to choose from. For each example, we can:

- **Counting:** Do all predictions then do a majority vote to decide class label
- **Knockout:** Compare two classes at a time, if one loses, never consider it again. Repeat till only one class remains.

What are the overall evaluation time complexities per example for Counting and Knockout, respectively?