

CM146, Fall 2019

Problem Set 4: Learning Theory, Boosting, Multi-class
Classification

Due December 8, 2019, 11:59pm

Name: Yongqian Li
Student ID #: 004997466

Submission instructions

- Submit your solutions electronically on the course Gradescope site as PDF files.
- If you plan to typeset your solutions, please use the LaTeX solution template. If you must submit scanned handwritten solutions, please use a black pen on blank white paper and a high-quality scanner app.

1 PAC Learning ^{all examples} [25 pts]

(a) Iterate in S_{train} , for each data point (x_i, y_i) in the m examples

• If its label $y_i = 1$, the hypothesis is h_{x_i} . ~~If none of~~
~~the~~ If after ^{all} iterations no $y_i = 1$ is found, the hypothesis
is h^- . The training error of this algorithm is always 0 since
the prediction is consistent with ~~the~~ S_{train} .

(b) The hypothesis $A(S_{\text{train}})$ is always consistent with S_{train} ,
and it makes at most 1 mistake in D since any $h^* \in H_{\text{singleton}}$
is either h^- or has only one $x \in X$ s.t. $h^*(x) = 1$, and mistakes are only made when
 $A(S_{\text{train}})$ is h^- with $h^* \neq h^-$ (Otherwise we would have predicted h^* correctly).

\therefore Loss is either 0 or $\frac{1}{N}$.

If the true h^* is h_{x^*} , we will make mistakes only if x^*
does not appear in S_{train} . $\therefore P(\text{loss} > 0) = (1 - \frac{1}{N})^m$

$\therefore P(\text{loss} > 0) \geq P(\text{loss} > \epsilon) \therefore$ If $\sigma \geq P(\text{loss} > 0)$, then $\sigma \geq P(\text{loss} > \epsilon)$

When $\text{loss} > \epsilon$, $\text{loss} = \frac{1}{N} \Rightarrow \frac{1}{N} > \epsilon \therefore (1 - \frac{1}{N})^m < (1 - \epsilon)^m$

\therefore let $\sigma \geq (1 - \epsilon)^m > (1 - \frac{1}{N})^m \geq P(\text{loss} > \epsilon)$,

$$\sigma \geq (1 - \epsilon)^m \Rightarrow e^{-\epsilon m} \leq \sigma$$

$$\text{if } -\epsilon m \leq \log \sigma$$

$$m \geq -\frac{\log \sigma}{\epsilon} = \frac{\log \frac{1}{\sigma}}{\epsilon}$$

$$\log \frac{1}{\sigma} \geq \frac{\log \frac{1}{\sigma}}{\epsilon}$$

$$\Rightarrow \sigma \geq (1 - \epsilon)^m$$

$$P(A(S_{\text{train}}) \neq h^*) \geq \epsilon$$

$$\therefore \text{If } m \geq \frac{\log \frac{1}{\sigma}}{\varepsilon},$$

$$\sigma \geq (1 - \varepsilon)^m > (1 - \frac{1}{N})^m$$

$$\therefore \text{Loss is either 0 or } \frac{1}{N} \therefore P(\text{loss} > \varepsilon) = P(\text{loss} > 0) = (1 - \frac{1}{N})^m$$

$$\therefore \sigma \geq (1 - \varepsilon)^m > P(\text{loss} > \varepsilon)$$

$$\therefore \text{when } m \geq \frac{\log \frac{1}{\sigma}}{\varepsilon}, \quad \sigma \geq P(\text{loss} > \varepsilon)$$

| i | Label | Hypothesis 1 (1st iteration) | | | | Hypothesis 2 (2nd iteration) | | | |
|-----|-------|------------------------------|----------------------|----------------------|----------------------|------------------------------|----------------------|-----------------------|-----------------------|
| | | D_0 | $f_1 \equiv [x > 2]$ | $f_2 \equiv [y > 6]$ | $h_1 \equiv [x > 2]$ | D_1 | $f_1 \equiv [x > 9]$ | $f_2 \equiv [y > 11]$ | $h_2 \equiv [y > 11]$ |
| (1) | (2) | (3) | (4) | (5) | (6) | (7) | (8) | (9) | (10) |
| 1 | - | 0.1 | - | + | - | 0.0625 | - | - | - |
| 2 | - | 0.1 | - | - | - | 0.0625 | - | - | - |
| 3 | + | 0.1 | + | + | + | 0.0625 | - | - | - |
| 4 | - | 0.1 | - | - | - | 0.0625 | - | - | - |
| 5 | - | 0.1 | - | + | - | 0.0625 | - | + | + |
| 6 | - | 0.1 | + | + | + | 0.25 | - | - | - |
| 7 | + | 0.1 | + | + | + | 0.0625 | + | - | - |
| 8 | - | 0.1 | - | - | - | 0.0625 | - | - | - |
| 9 | + | 0.1 | - | + | - | 0.25 | - | + | + |
| 10 | + | 0.1 | + | + | + | 0.0625 | - | - | - |

Table 1: Table for Boosting results

2 VC Dimension [15 pts] *directions*

$VC(H) = 3$. $ax^2 + bx + c$ has two ~~graphs~~ of graphs on a Cartesian plane.

for any point:

for 2 points:

$\therefore VC(H) \geq 2$

for 3 points:

$\therefore VC(H) \geq 3$

for 4 points, This assignment cannot be shattered by H.

$\therefore VC(H) = 3$

3 Boosting [40 pts]

| i | x | y | Label |
|----|----|----|-------|
| 1 | 0 | 8 | - |
| 2 | 1 | 4 | - |
| 3 | 3 | 7 | + |
| 4 | -2 | 1 | - |
| 5 | -1 | 13 | - |
| 6 | 9 | 11 | - |
| 7 | 12 | 7 | + |
| 8 | -7 | -1 | - |
| 9 | -3 | 12 | + |
| 10 | 5 | 9 | + |

(a). $D_0(i) = \frac{1}{10} = 0.1$ $\therefore VC(H) = 3$

(b). $\epsilon_{f_1} = \frac{2}{10} = 0.2$
 $\epsilon_{f_2} = \frac{3}{10} = 0.3 > \epsilon_{f_1}$
 \therefore choose f_1

(c). $\epsilon_1 = 0.2$
 $\alpha_1 = \frac{1}{2} \log_2 \left(\frac{1 - 0.2}{0.2} \right)$
 $= \frac{1}{2} \log_2 \left(\frac{0.8}{0.2} \right)$
 $= \frac{1}{2} \log_2 4$

$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \cdot \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) \neq y_i \\ e^{\alpha_t} & \text{else} \end{cases}$

$Z_1 = 2 \sqrt{\epsilon_1 (1 - \epsilon_1)}$
 $= 2 \sqrt{0.2 \times 0.8} = 0.8$

$$\begin{aligned} \therefore D_0(i) &= 0.1 \\ \therefore \text{when } h_1(x_i) = y_i, D_1(i) &= \frac{0.1 \cdot 2^{-1}}{0.8} = \frac{0.1}{0.8 \times 2} \approx 0.0625 \\ \text{when } h_1(x_i) \neq y_i, D_1(i) &= \frac{0.1 \cdot 2^1}{0.8} \approx 0.25 \end{aligned}$$

(d). ~~$H_{\text{final}} = \text{sgn}(\sum \alpha_t h_t(x))$~~ $\epsilon_2 = 0.0625 \times 4 = 0.25$
 $\alpha_2 = \frac{1}{2} \log_2 \left(\frac{1 - 0.25}{0.25} \right) = \frac{1}{2} \log_2 3 \approx 0.792$

$$\begin{aligned} H_{\text{final}}(x) &= \text{sgn} \left(\sum_t \alpha_t h_t(x) \right) \\ &= \text{sgn} (h_1(x) + 0.792 h_2(x)) \\ &= \text{sgn} ([x > 2] + 0.792 [y > 11]) \end{aligned}$$

4 Multi-class classification [60 pts]

Consider a multi-class classification problem with k class labels $\{1, 2, \dots, k\}$. Assume that we are given m examples, labeled with one of the k class labels. Assume, for simplicity, that we have m/k examples of each type.

Assume that you have a learning algorithm L that can be used to learn Boolean functions (E.g., think about L as the Perceptron algorithm). We would like to explore several ways to develop learning algorithms for the multi-class classification problem.

There are two schemes to use the algorithm L on the given data set, and produce a multi-class classification:

- **One vs. All:** For every label $i \in [1, k]$, a classifier is learned over the following data set: the examples labeled with the label i are considered "positive", and examples labeled with any other class $j \in [1, k], j \neq i$ are considered "negative".
- **All vs. All:** For every pair of labels (i, j) , a classifier is learned over the following data set: the examples labeled with one class $i \in [1, k]$ are considered "positive", and those labeled with the other class $j \in [1, k], j \neq i$ are considered "negative".

- (a). i. for One vs. All, k classifiers are learned;
 for All vs. All, $\frac{k(k-1)}{2}$ classifiers are learned.
- ii. for One vs. All, m examples;
 for All vs. All, $\frac{2m}{k}$ examples.
- iii. for One vs. All, the label $y_i = \underset{y \in \{1, 2, \dots, k\}}{\text{argmax}} w_y^T x_i$
 for All vs. All, each label gets $(k-1)$ votes,
 there are 2 methods:

(1). Tournament: start with any pair of labels $\langle a, b \rangle$, compare $w_j^T x$ and $w_i^T x$ and $w_b^T x$ and $w_a^T x$; the winner (with larger $w^T x$) compare with another label; continue with winners until the final winner is found

(2). Majority: compare each pair of labels $\langle a, b \rangle$, the larger one of $w_a^T x$ and $w_b^T x$ indicates the winner. Then choose the label which wins more often than any other label.

~~I would choose All vs. All because for One against All, each class needs to be separable from the others, which is a very strong assumption. Each comparison of One vs. All involves a class with k examples and another class with k examples, so they may be very unbalanced.~~

iv. The computational complexity for One vs. All is $k \cdot m \Rightarrow O(k \cdot m)$; that of All vs. All is $\frac{k(k-1)}{2} \cdot \frac{m}{k} \Rightarrow O(k \cdot m)$. ~~d is the dimension of each example.~~

(b). I would choose ~~All~~ ^{One} vs. All because it is ~~more efficient~~ and only need to store k classifiers instead of $\frac{k(k-1)}{2}$. It is also easier to implement, and performs as good as ~~All~~ vs. All in this case.

(c). ~~Let d be the dimension of each example.~~ Kernel Perceptron calculates ~~the~~ ^{each} classifier in $O(\frac{m^2}{n^2})$. \therefore One vs. All: $O(km^2)$
All vs. All: $\frac{k(k-1)}{2} \cdot \frac{m}{k} \Rightarrow O(m^2)$
 \therefore Now ~~All~~ vs. All is more efficient, so I ~~for~~ would prefer ~~All~~ vs. All.
~~since it needs $K(x_i, x_j)$ between each pair~~

(d). One vs. All: $k \cdot O(dm^2) \Rightarrow O(kdm^2)$
All vs. All: $\frac{k(k-1)}{2} \cdot O(\frac{dm^2}{k}) \Rightarrow O(dm^2)$
 \therefore All vs. All is ~~more~~ more efficient.

(e). One vs. All: $k \cdot O(d^2m) \Rightarrow O(kd^2m)$
All vs. All: $\frac{k(k-1)}{2} \cdot O(\frac{d^2m}{k}) \Rightarrow O(kd^2m)$
 \therefore They are the same efficient.

(f). Counting: for each example, test all $\frac{k(k-1)}{2}$ classifiers
 $\therefore \frac{k(k-1)}{2} \cdot d \Rightarrow O(k^2d)$

Knockout: for each example, eliminate 1 classifier at each step until the winner is found. $\therefore (k-1)$ steps.
 $(k-1) \cdot d \Rightarrow O(kd)$